

## SUMMARY

Senior Backend Engineer with 6 years of experience building highly available, scalable cloud systems processing millions of daily transactions. A collaborative builder and critical thinker who values simplicity and low-risk solutions. Led architecture design for 10+ technical designs across microservices, Machine Learning systems, and real-time pipelines. Combines an innovative, highly adaptive mindset with deep engineering accountability. Active technical blogger at cuongkane.com.

## EXPERIENCE

- **Parcel Perform** Ho Chi Minh City, Vietnam  
Aug 2024 - Present
  - *Senior Backend Engineer*
    - **System Architecture & Design:** Directed architecture for 10+ systems including ML and real-time pipelines; authored cross-team Technical Design for carrier identification spanning 7 microservices, approved by CTO.
    - **Microservices & Database Splitting:** Proposed and Engineered decomposition of 2 monoliths into microservices with zero-downtime database splitting, enhancing modularity and enabling autonomous team deployments, improving team velocity by 30%.
    - **High-Availability Cloud Systems:** Proactively take over and maintain mission-critical analytics platform (ClickHouse/Dagster); accelerated dashboard chart loading by implementing static file caching strategies.
    - **Real-Time Data Pipelines:** Optimized ML-based carrier identification pipeline processing millions of daily events across 7 microservices; increased success rate by seamlessly integrating ML models with legacy regex flows.
    - **AI/ML Solutions:** Pioneered building AI agents using LangChain following ReAct architecture and AWS Bedrock; reduced LLM costs by 20% via strategic model selection while automating Root Cause Analysis workflows.
    - **Engineering Excellence:** Pioneered company-wide adoption of AI-assisted Technical Design workflows, ensuring code quality with LLM.
    - **Release & Deployment Strategy:** Achieved zero-downtime releases by establishing canary deployment protocols, A/B testing, and mandatory QA gates for all production changes.
    - **Observability & Sustainable Growth:** Enhanced system reliability by defining and monitoring SLA/SLO targets; utilized percentile latencies and Kafka lag metrics to drive proactive capacity planning and ensure consistent service availability.
    - **Mentorship & Knowledge Sharing:** Onboarded 5+ engineers, mentored 1 engineer to promotion-readiness and fostered learning culture by delivering 15+ technical sessions to entire dev teams on System Design, AI/LangChain, and software engineering best practices.
    - **Engineering Leadership:** Strengthened engineering standards by sharing culture, showcasing PEP8/TDD best practices, and conducting 6+ technical interviews.
- **Parcel Perform** Ho Chi Minh City, Vietnam  
Jul 2022 - Aug 2024
  - *Backend Engineer*
    - **Engineering Excellence:** Maintained stable and fast development velocity (compared to other devs) by applying best practices (TDD) and effective task management.
    - **Scalable API Development:** Designed and Launched Public Booking API for multi-carrier integrations; and Celery-based rate limiting to protect third-party resources.
    - **Extensible Integration Architecture:** Designed extensible integration framework supporting 50+ carriers; validated requirements through usage testing to ensure scalability for future logistics partners.
    - **High-Volume Data Pipelines:** Scaled data pipeline to process 5M+ daily events for mission-critical tracking using Apache Flink and optimized microservices.
    - **Email Extraction Platform:** Engineered high-throughput email extraction engine using Kafka CDC; system has processed shipment notifications reliably in production for 3+ years.
    - **Sync-to-Async Reliability Pattern:** Ensured booking reliability by architecting a sync-to-async fallback pattern across 5 microservices using Kafka to handle component failures.
    - **Celery Pipeline Architecture:** Accelerated data processing by parallelizing 8K+ postal code tasks using optimized Celery canvas patterns (chain/chord) and worker configurations.

- **System Reliability & Incident Response:** Restored 1.5M records during critical incidents within 24 hours by developing custom recovery scripts and resolving Redis cluster hotspots under 10x traffic load.
- **Large-Scale Data Migrations:** Secured data consistency during large-scale migrations by implementing incremental update patterns, optimizing SQL indexing, and managing autovacuum constraints.
- **Testing & Developer Experience:** Defined and spread company-wide testing standards (AAA pattern) and drove adoption of modern tooling (Poetry, Ruff).

### Cinnamon AI Labs

Ho Chi Minh City, Vietnam

Aug 2020 - Jul 2022

- *Software Engineer*

- **Enterprise Client Delivery:** Deployed enterprise-grade document processing systems for large enterprise clients (Deloitte, Bridgestone,...), directly driving client success in AI automation.
- **Multi-Framework Backend Development:** Developed versatile backend solutions using Flask, FastAPI, and Django, strictly adhering to SOLID principles and Repository patterns.
- **Multi-Cloud Infrastructure:** Managed multi-cloud infrastructure (AWS, Azure, GCP) and containerized applications with Docker to ensure consistent cross-cloud deployments.
- **ML-to-Production Pipeline:** Bridged gap between research and production by converting experimental ML models into scalable, SLA-compliant API services.

## PROJECTS

---

- **Codlis - AI Code Visualization (Master's Thesis, H2 2025) ([codlis.com](http://codlis.com))**: Architected multi-agent LLM system for interactive code comprehension using LangChain and Knowledge Graphs. Validated effectiveness via user study showing 2.9x higher task success rate and 29% cognitive load reduction.
- **SweatCharge - Sports Club Management (H1 2025) ([sweatcharge.com](http://sweatcharge.com))**: Developed full-stack sports management platform with role-based SSO and automated billing; currently servicing active production users. Implemented fuzzy matching algorithms for player association using React/TypeScript.
- **Hasiti - AI Hairstyle Simulator (Bachelor's Thesis, 2021)**: Built real-time hairstyle simulation app using fine-tuned CNNs for style transfer. Architected full-stack solution connecting React Native mobile client with Python/Flask AI inference backend.

## CERTIFICATIONS

---

- **AWS Certified Solutions Architect - Associate**: Amazon Web Services (Valid: Jan 2025 – Jan 2028)

## EDUCATION

---

- **Master of Science in Computer Science**  
*HCM University of Science*

Sept. 2021 – Sept. 2025  
Ho Chi Minh City, Vietnam

- **Research Focus:** AI/ML, Big Data, and Human-Computer Interaction under post-doctoral supervision. Applied analytical experimentation methods for metric-driven evaluation.

- **Honor Program - Bachelor of Science in Information Technology**  
*HCM University of Science*

Sept. 2017 – Sept. 2021  
Ho Chi Minh City, Vietnam

- **Honor Program:** Top 70 students selected from 1,000+ candidates. Comprehensive curriculum spanning web/mobile development, computer vision, image processing, software engineering, and scientific research methods.
- **Technical Foundation:** Developed fundamentals in algorithms, data structures, OOP principles, and hardware/software architecture through rigorous coursework.

- **Machine Learning Course (Facebook-sponsored)**  
*Coder School*

Aug. 2019 – Dec. 2019  
Ho Chi Minh City, Vietnam

## AWARDS & RECOGNITION

---

- **Parcel Perform (2022-Present)**: Member of the Month (3x), Best Presentation Award, AVA AI Award for AI adoption, Professional Style Award.
- **Cinnamon AI Labs (2020-2022)**: Employee Recognition Award (Monthly) for outstanding technical contributions.
- **Faculty of IT Academic Excellence (2018)**: Top 10 students with highest GPA (9.08/10) among 1,000+ students.

## TECHNICAL SKILLS

---

- **Architecture:** High-Availability Systems, Microservices, System Design, RESTful APIs, TDD, Idempotency, Outbox Pattern, Feature Toggles, SSO/OAuth
- **Big Data & Pipeline:** Apache Flink, Kafka, Celery (canvas: chain/chord/group), Dagster, RabbitMQ
- **Infrastructure & Observability:** AWS (EC2, S3, Bedrock), Docker, Grafana, Sentry, Alerting Systems, CI/CD
- **AI/ML:** LangChain, LangGraph, LangSmith, AI Agents, Prompt Engineering
- **Languages:** Python, Scala, Java, TypeScript, SQL
- **Databases & Data Stores:** PostgreSQL, Redis, ClickHouse
- **Web Frameworks:** Django, Django Ninja, FastAPI, Pydantic, Next.js/React (side projects)
- **Engineering Practices:** Technical Planning, Task Management, Story Estimation, Sprint Planning, Agile
- **Dev Environment:** Vim (primary editor), tmux, terminal-centric workflow, split keyboard enthusiast

## LANGUAGES

---

- **English:** B2 Upper Intermediate – Aptis ESOL (British Council) 176/200 **Vietnamese:** Native